

Machine Learning 1

Lecture 10.1 - Unsupervised Learning
Principal Component Analysis - Variance
Maximization

Erik Bekkers

(Bishop 12.1.1)



Continuous latent space

Goal

- Dimensionality reduction: model the data in a low dim. space
- Example: take one grey-scale image of “3” and make multiple copies by translation and rotation

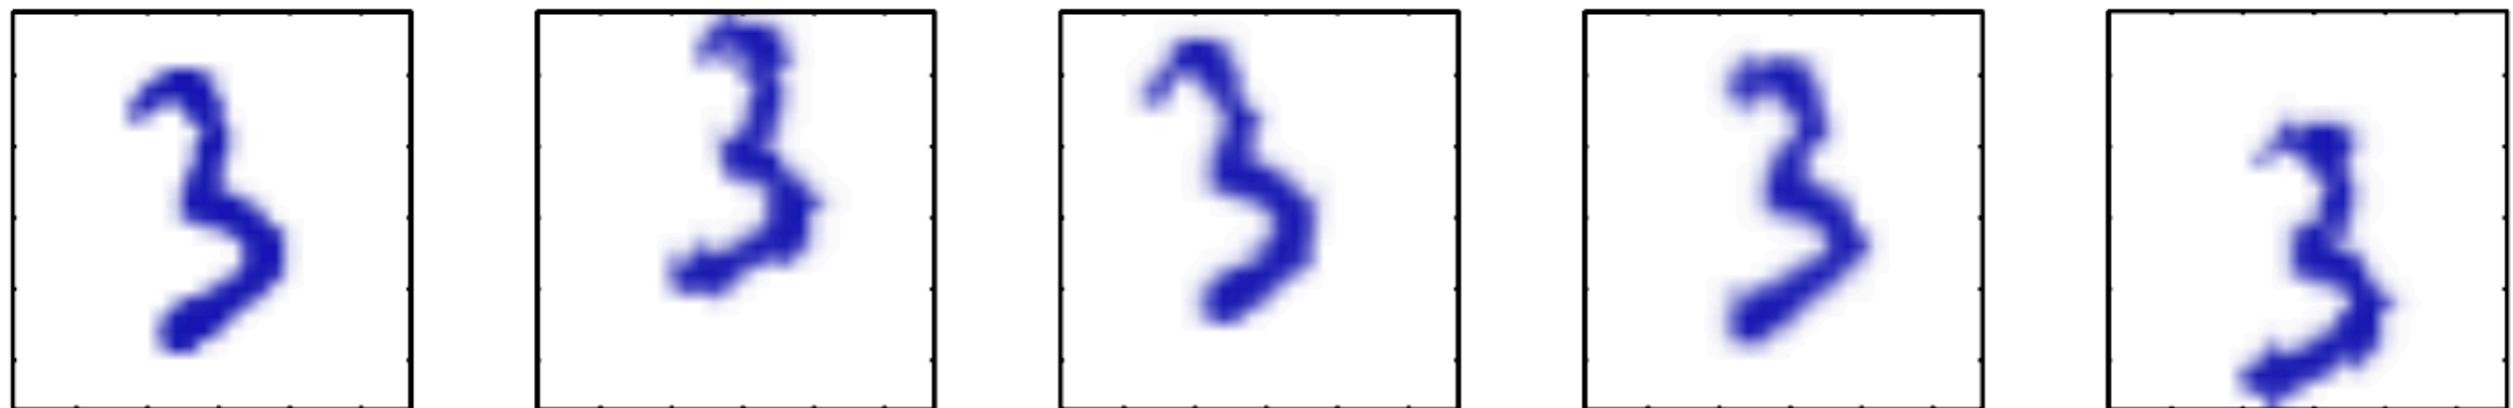


Figure: Synthetic “3” dataset (Bishop 12.1)

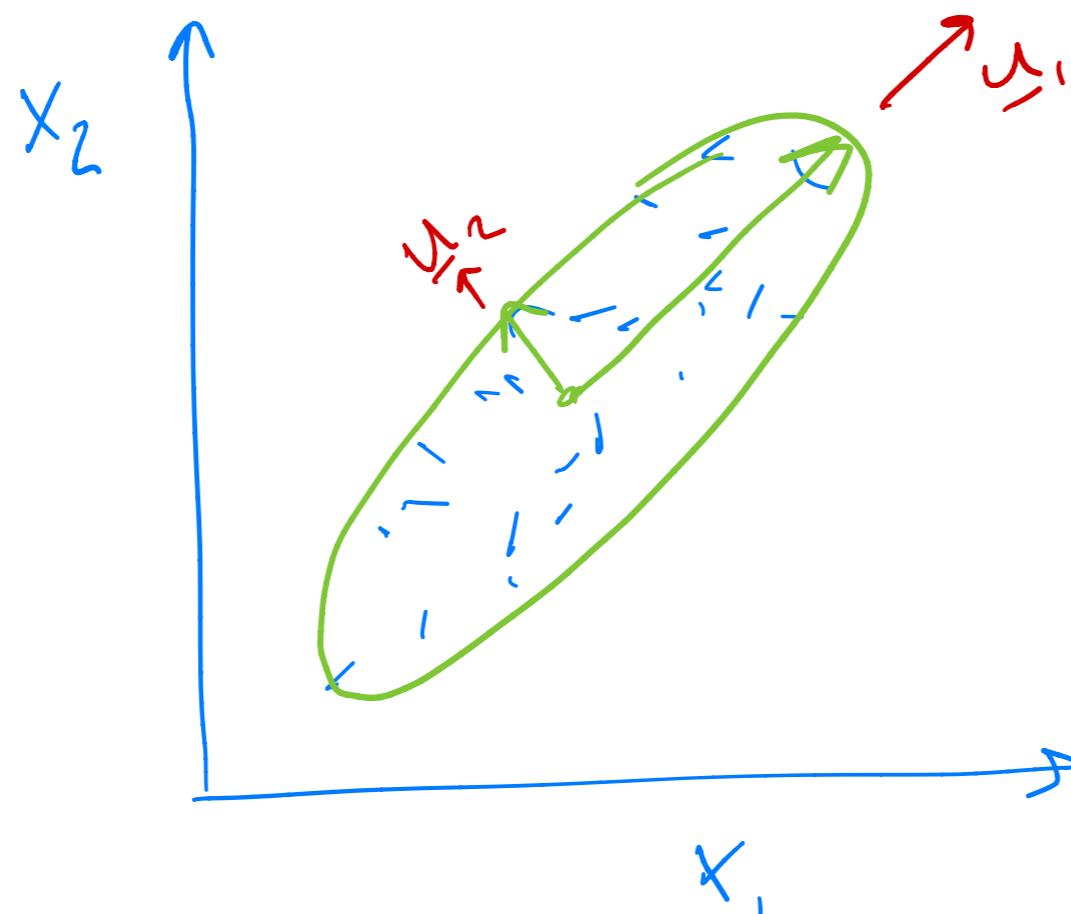
- Pixel space dimension: 100x100 pixels
- Latent space dimension: $3 = 2$ (translations) + 1 rotation
- From the 3 latent variables we could generate all 100x100 pixels!

Example continued

- › A more realistic dataset of images will have more degrees of freedom in the latent space, such as:
 - › Scaling
 - › Digits from 0-9
 - › Colors
 - › Different hand-writing styles
 - › Etc.
- ... but still much fewer than 100x100!
- › In this example, the latent subspace is a non-linear transformation of the images
- › We first study linear latent spaces with PCA and later consider generalizations to the non-linear case

Principal Component Analysis (PCA)

- ▶ Find a linear projection of the data such that the variance in the projected space is maximal
- ▶ PCA captures the axes of maximal variation in the data, called **principal components**



Principal Component Analysis (PCA)

- › Data: $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}, \mathbf{x}_n \in \mathbb{R}^D$
- › Goal: project data into a $M < D$ dimensional space while **maximizing the variance** of the projected data
- › M is given
- › Mean and covariance defined by

$$\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

$$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T$$

- › \mathbf{S} is symmetric and positive definite




Figure: Maximizing variance of projections (Bishop 12.2)

1D Projection

- Project data int the first latent dimension by a vector $\mathbf{u}_1 \in \mathbb{R}^D$
 $\underbrace{\mathbf{x}_n}_{z_n \in \mathbb{R}}$
- The projection gives the scalar $\mathbf{u}_1^T \mathbf{x}_n$, the mean of the projection is $\mathbf{u}_1^T \bar{\mathbf{x}}$
- We only need its direction, so normalize this component: $\|\mathbf{u}_1\|^2 = \mathbf{u}_1^T \mathbf{u}_1 = 1$
- The variance of the projected data is

$$\begin{aligned}\text{var}[z_i] &= \frac{1}{N} \sum_{n=1}^N (\mathbf{u}_1^T \mathbf{x}_n - \mathbf{u}_1^T \bar{\mathbf{x}})^2 = \frac{1}{N} \sum_{n=1}^N (\mathbf{u}_1^T (\mathbf{x}_n - \bar{\mathbf{x}}))^2 \\ &= \frac{1}{N} \sum_{n=1}^N \mathbf{u}_1^T (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T \mathbf{u}_1 \\ &= \mathbf{u}_1^T \left(\frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T \right) \mathbf{u}_1 = \boxed{\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1}\end{aligned}$$

$\underbrace{\text{cov}[\mathbf{x}]}_{\mathbf{S}} = \mathbf{S}$

$E[z] = E[\mathbf{u}_1^T \mathbf{x}]$
 $= \mathbf{u}_1^T E[\mathbf{x}]$

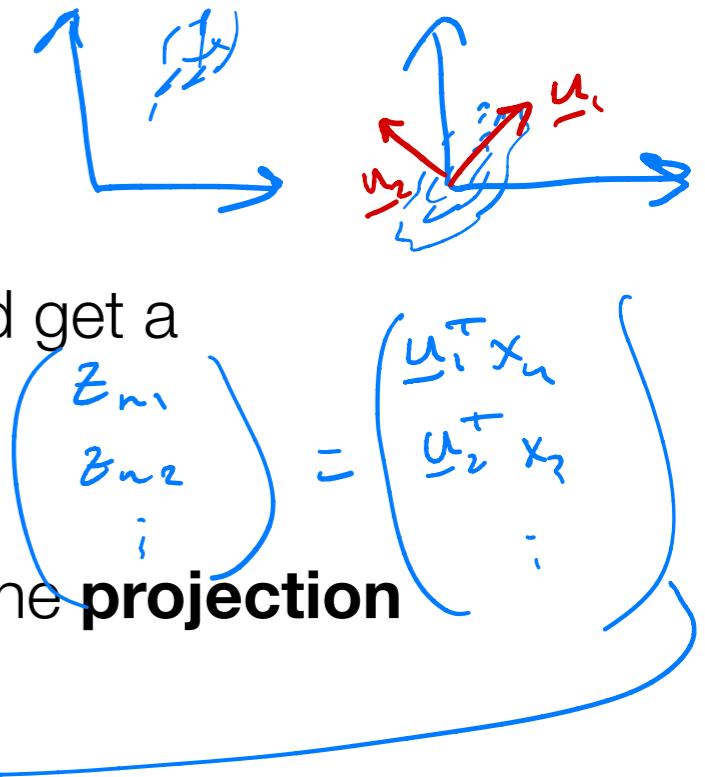
Maximizing the variance of 1 component

Need this constraint
else objective is infinite

- ▶ Solve $\underset{\mathbf{u}_1}{\operatorname{argmax}} \underbrace{\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1}_{f(\mathbf{u}_1)}$ subject to $\underbrace{\mathbf{u}_1^T \mathbf{u}_1 = 1}_{g(\mathbf{u}_1) = c}$
- ▶ Method of Lagrange multipliers
 - ▶ Define Lagrangian $L(\mathbf{u}_1, \lambda_1) = \underbrace{\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1}_{f(\mathbf{u}_1)} + \lambda_1 (\mathbf{u}_1^T \mathbf{u}_1 - 1)$
 - ▶ Solving for \mathbf{u}_1 means $\frac{\partial}{\partial \mathbf{u}_1} L(\mathbf{u}_1, \lambda_1) = \mathbf{S} \mathbf{u}_1 - \lambda_1 \mathbf{u}_1 = 0$
 - ▶ We need to solve eigensystem $\boxed{\mathbf{S} \mathbf{u}_1 = \lambda_1 \mathbf{u}_1}$
 - ▶ So \mathbf{u}_1 and λ_1 are respectively an eigenvector and eigenvalue of $\mathbf{S} \in \mathbb{R}^{D \times D}$!
- ▶ The \mathbf{u}_1 is called a **principal component.** $\mathbf{u}_1^T \mathbf{A} \mathbf{u}_1 = \lambda_1 \mathbf{u}_1^T \mathbf{u}_1 = \lambda_1$
- ▶ The variance of the projected data is $\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 = \lambda_1$
- ▶ Maximizing variance means we search for the eigenvector with **largest eigenvalue**

PCA via maximum variance

- We repeat the procedure for M orthogonal vectors and get a projection defined by $U_M = [\mathbf{u}_1, \dots, \mathbf{u}_M] \in \mathbf{R}^{D \times M}$
- PCA: compute $\bar{\mathbf{x}}$ and the eigen-decomposition of \mathbf{S} . The **projection** then is $\mathbf{z} = U_M^T(\mathbf{x} - \bar{\mathbf{x}})$
- Those are M eigenvectors of \mathbf{S} , **the principal components**. The eigenvalues are $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_M$
- The matrix \mathbf{S} is positive semi-definite, thus $\forall_j : \lambda_j \geq 0$
- The (total) variance of the projected data is $\text{Tr}[\text{Cov}[\mathbf{z}]] = \sum_{j=1}^M \lambda_j$



Reminder: eigen-decomposition

- When the matrix is **symmetric positive semi-definite**:


$$\underline{u}_i^\top \underline{u}_j = \begin{cases} 1 & \text{if } i=j \\ 0 & \text{otherwise} \end{cases}$$

$$\mathbf{S} = \mathbf{U} \Lambda \mathbf{U}^T \quad \text{with} \quad \Lambda = \text{diag}\{\lambda_1, \dots, \lambda_D\}$$

- The eigenvectors are **orthonormal** and are stored in $\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_D]$
Change of basis → rotation
- All eigenvalues are **non-negative** and are the elements of the diagonal matrix Λ

$$= \text{Tr}(\Lambda \mathbf{U}^T \mathbf{U}) = \text{Tr}(\Lambda I)$$

- Total variance given by $\text{Tr}(\mathbf{S}) = \text{Tr}(\mathbf{U} \Lambda \mathbf{U}^T) = \text{Tr}(\Lambda) = \sum_{i=1}^D \lambda_i$



Getting the eigenvectors in practice

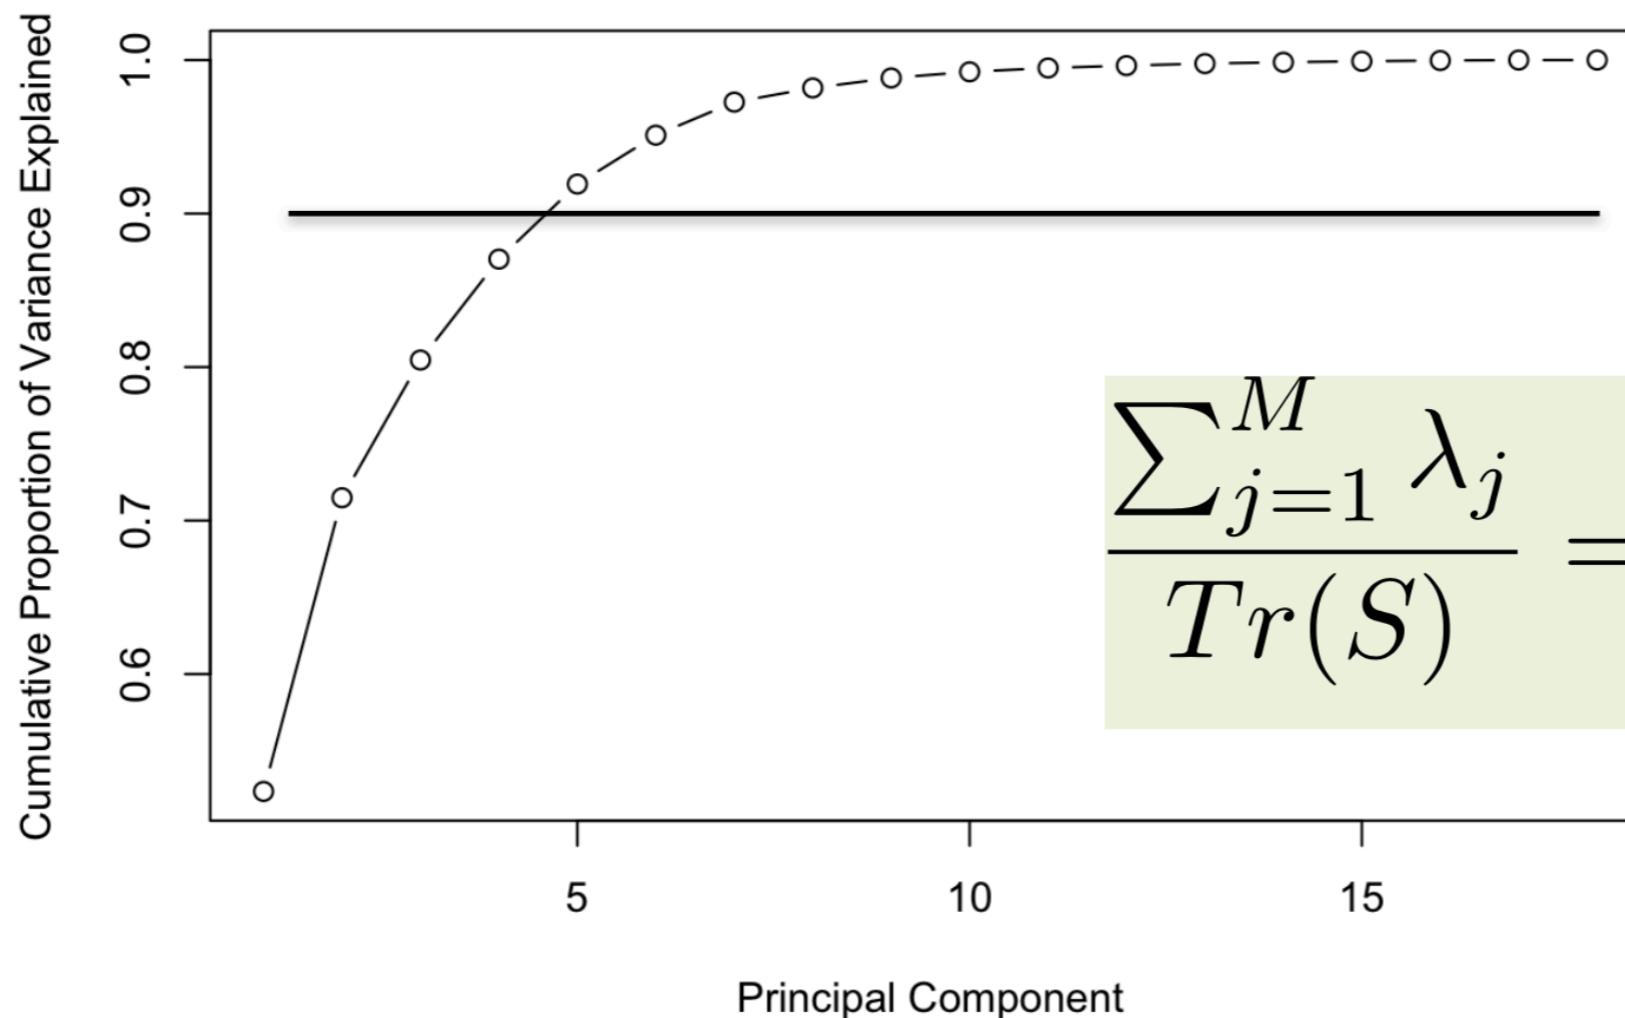
- ▶ Full eigenvalue decomposition is expensive: $O(D^3)$
- ▶ Only need up to the M^{th} component: $O(MD^2)$
- ▶ In python:

```
M = 10
S = np.cov(X)
Um, Lm, Vm = scipy.sparse.linalg.svds(S, k=M)
```

For symmetric positive definite matrices such as S , the SVD decomposition is equivalent to the eigen-decomposition

How to choose M ?

- › We can measure the discarded variance
- › For example to preserve 90% of the variance, pick M such that



The proportion of explained variance

$$\frac{\sum_{j=1}^M \lambda_j}{Tr(S)} = \frac{\sum_{j=1}^M \lambda_j}{\sum_{i=1}^D \lambda_i} > 0.9$$



Applications: dimensionality reduction

- › When data is defined in high dimension (large D) we want to project down to lower dimension because:
 - › Reduce time and storage space required
 - › For classification/regression: our model **will have less parameters**, thus we need less data points for learning
- › Other methods (not covered): **feature selection**. PCA is known as a **feature extraction** method instead.

Applications: 2D Visualization (MNIST)



MNIST: 24 x 24 pixels




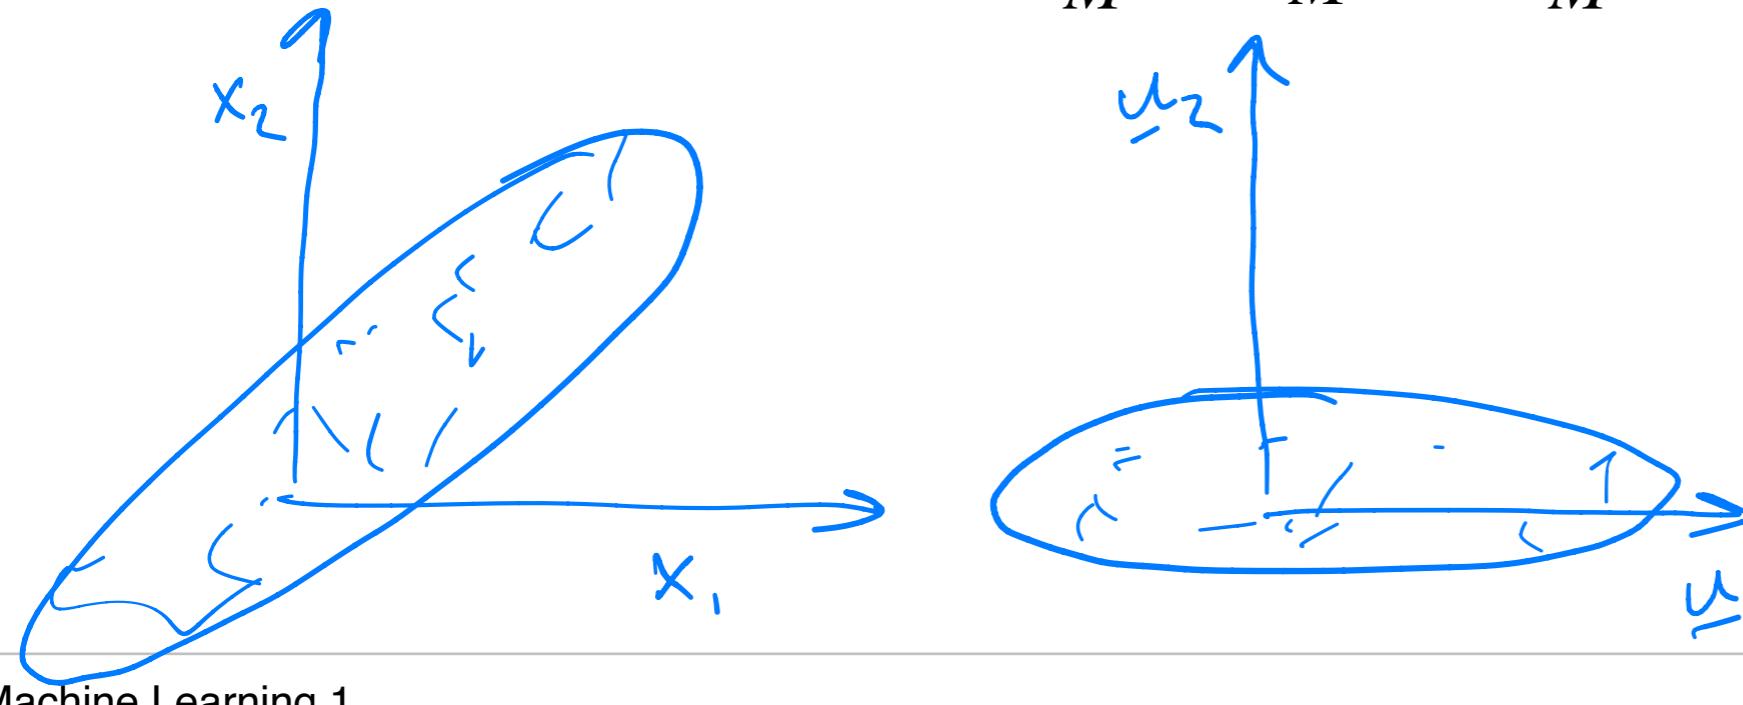
Feature Decorrelation

- Good side effect of PCA: features have **no correlation** in the projected space.
- The covariance matrix of the projected data is **diagonal**

$$\frac{1}{N} \sum_{n=1}^N \mathbf{z}_n \mathbf{z}_n^T = \frac{1}{N} \sum_{n=1}^N \mathbf{U}_M^T (\mathbf{x}_n - \bar{\mathbf{x}}) (\mathbf{x}_n - \bar{\mathbf{x}})^T \mathbf{U}_M$$

$\mathbf{u}_i^T S \mathbf{u}_i = \lambda_i$

$$= \mathbf{U}_M^T \mathbf{S} \mathbf{U}_M = \mathbf{U}_M^T \mathbf{U} \Lambda \mathbf{U}^T \mathbf{U}_M = \Lambda_M$$



Applications: whitening (or spherering)

- Before applying learning algorithms we usually do some pre-processing:
 - e.g. **standardization**: subtract the mean and divide by the standard deviation
- With PCA we can **whiten** the data, one step more:
 - Centre** and **de-correlate** the features:

$$\mathbf{z} = \mathbf{U}_M^T(\mathbf{x} - \bar{\mathbf{x}})$$

- Cast features to **unit standard deviation** by rescaling:

$$\mathbf{z} = \Lambda_M^{-1/2} \mathbf{U}_M^T(\mathbf{x} - \bar{\mathbf{x}})$$

